




Observability for Distributed Computing with Dask

How to remain sane while identifying
and solving your problems.

Hendrik Makait

@hendrikmakait

hendrik@coiled.io



dask



General parallelism for Python

- Parallel for-loopy code
- Dataframes / ETL
- Array computing
- Dynamic and complex systems / ML





General parallelism for Python

- Parallel for-loopy code
- Dataframes / ETL
- Array computing
- Dynamic and complex systems / ML



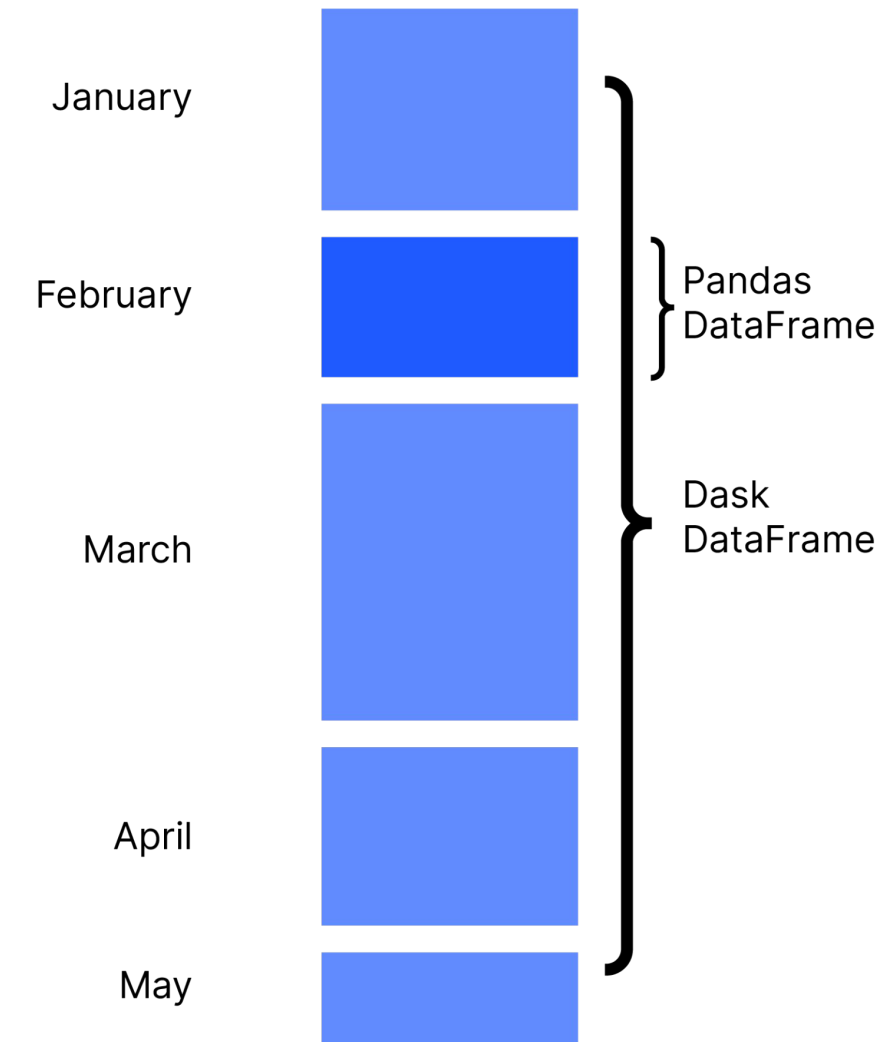
```
futures = []  
for filename in filenames:  
    future = client.submit(  
        process,  
        filename,  
    )  
    futures.append(future)
```





General parallelism for Python

- Parallel for-loopy code
- Dataframes / ETL
- Array computing
- Dynamic and complex systems / ML



```
import dask.dataframe as dd
```

```
df = dd.read_csv("s3://.../*.csv")
```

```
...
```


```
df.to_parquet("s3://.../clean.parquet")
```





General parallelism for Python

- Parallel for-loopy code
- Dataframes / ETL
- Array computing
- Dynamic and complex systems / ML



```
import dask.array as da
```

```
x = da.from_array(...)
```

```
x -= x.mean(axis=2)
```


```
x.dot(x.T).sum(mean=0)
```





General parallelism for Python

- Parallel for-loopy code
- Dataframes / ETL
- Array computing
- Dynamic and complex systems / ML



```
import prefect
import xgboost.dask
import lightgbm
import optuna
import joblib
import xarray
```

...



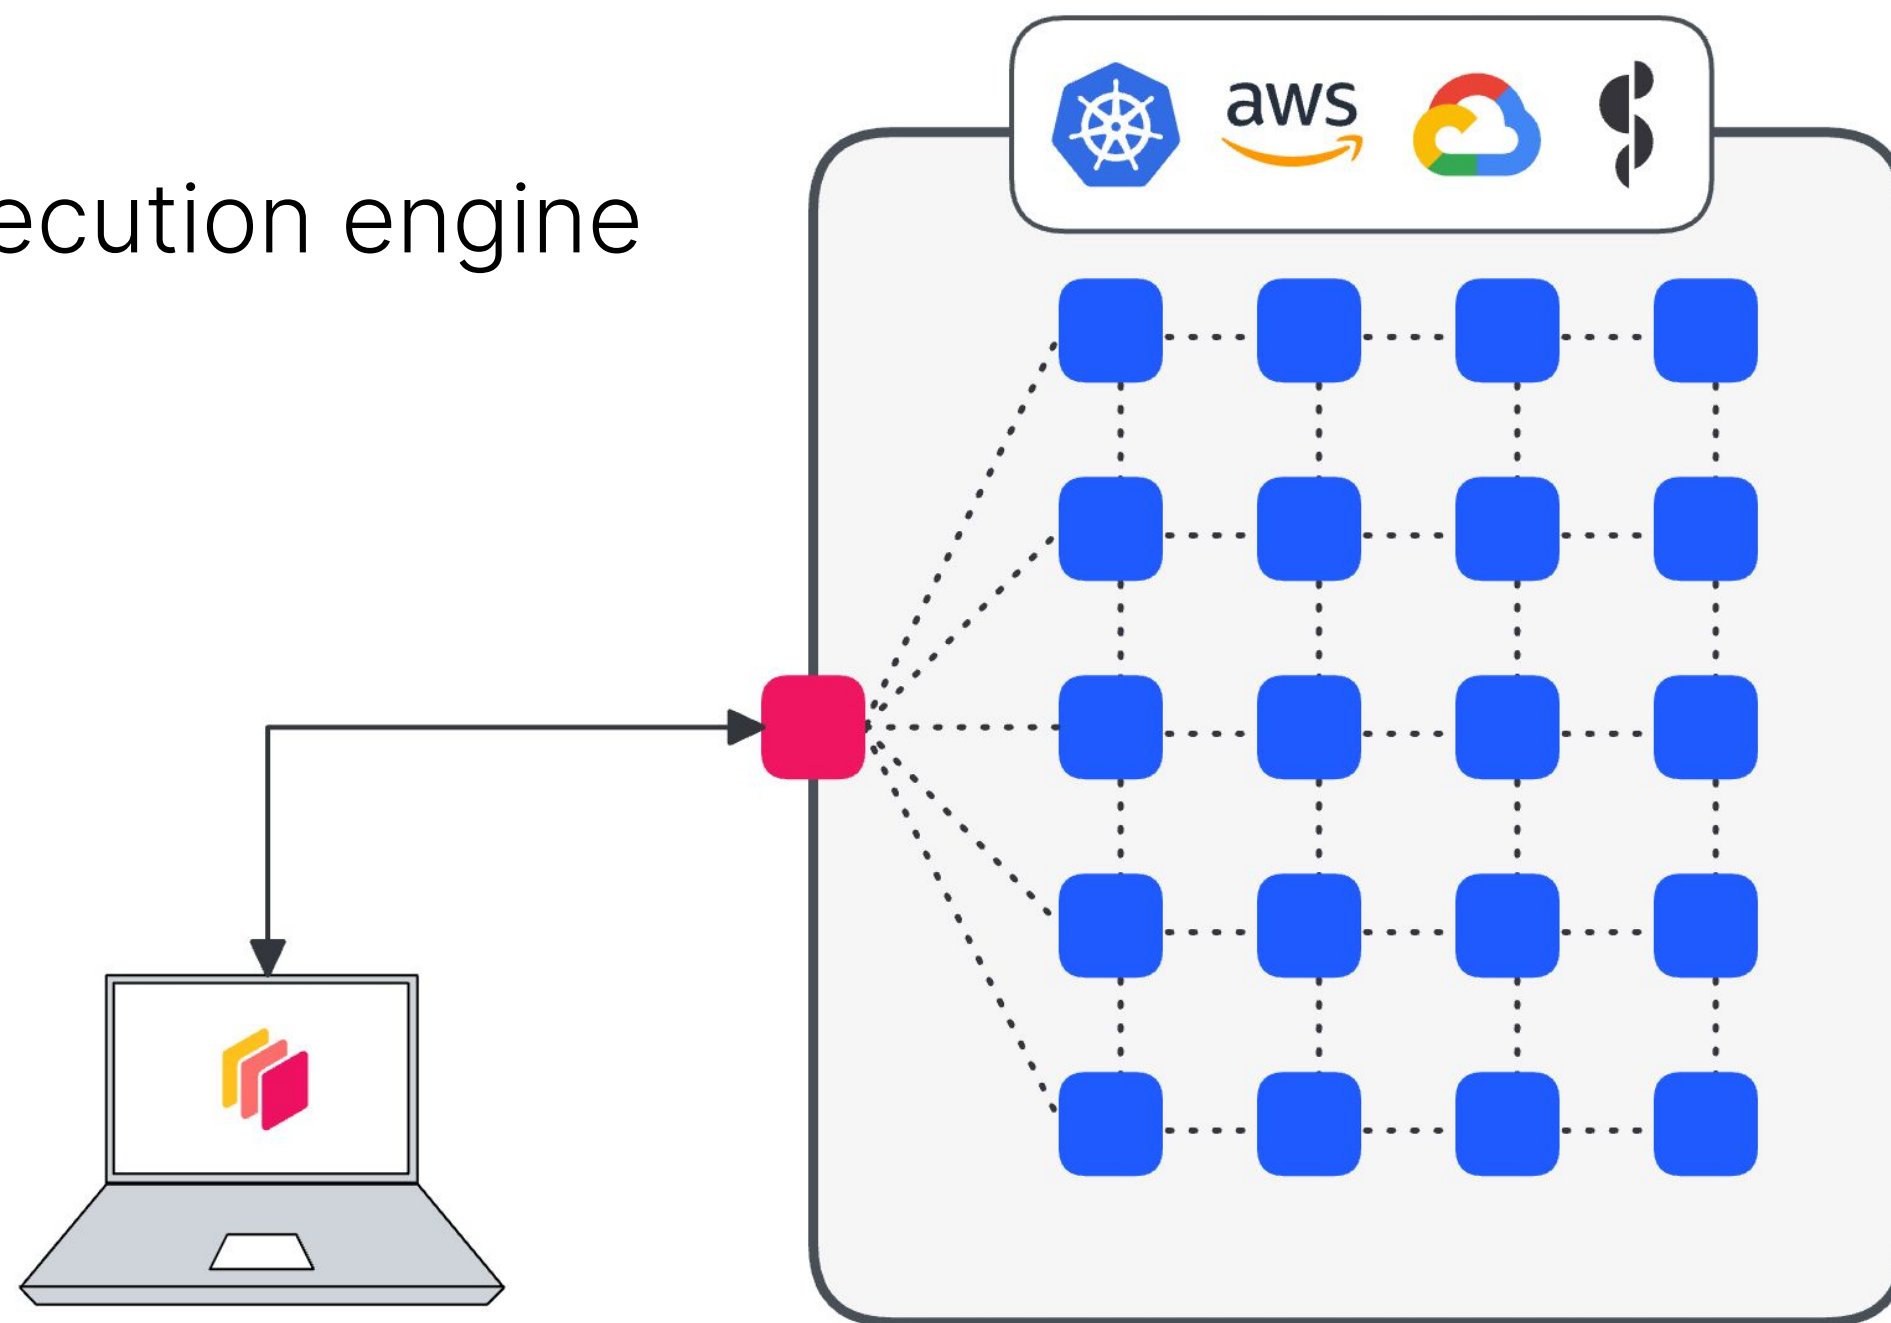


Distributed computing for Python

Scales Python with its distributed execution engine

Several deployment options:

- Kubernetes
- HPC
- Cloud
- Coiled





Observability



Observability

What is Observability?

“Observability lets us **understand a system from the outside**, by letting us ask questions about that system **without knowing its inner workings**. Furthermore, it allows us to easily troubleshoot and handle novel problems (i.e. “unknown unknowns”), and helps us answer the question, **“Why is this happening?”**”


– <https://opentelemetry.io/docs/concepts/observability-primer/>

→ An observable system enables us to monitor and debug it.



The pillars of observability

Logging



Metrics



Tracing



Diagnostic Tooling





Logging



```
2023-03-20 15:14:43,852 - distributed.nanny - INFO - Start Nanny at: 'tls://10.244.10.11:33815'
2023-03-20 15:14:46,216 - distributed.worker - INFO - Start worker at:  tls://10.244.10.11:46721
2023-03-20 15:14:46,216 - distributed.worker - INFO - Listening to:  tls://10.244.10.11:46721
2023-03-20 15:14:46,216 - distributed.worker - INFO - Worker name: dask-worker-351d74e2f7f644beb571a17408e81902-8z6lj
2023-03-20 15:14:46,216 - distributed.worker - INFO - dashboard at: 10.244.10.11:8787
2023-03-20 15:14:46,216 - distributed.worker - INFO - Waiting to connect to: tls://dask-351d74e2f7f644beb571a17408e81902.staging:8786
2023-03-20 15:14:46,216 - distributed.worker - INFO - -----
2023-03-20 15:14:46,216 - distributed.worker - INFO - Threads: 1
2023-03-20 15:14:46,217 - distributed.worker - INFO - Memory: 8.00 GiB
2023-03-20 15:14:46,217 - distributed.worker - INFO - Local Directory: /tmp/dask-worker-space/worker-4j9p24_r
2023-03-20 15:14:46,217 - distributed.worker - INFO - -----
2023-03-20 15:14:46,235 - distributed.worker - INFO - Registered to: tls://dask-351d74e2f7f644beb571a17408e81902.staging:8786
2023-03-20 15:14:46,235 - distributed.worker - INFO - -----
2023-03-20 15:14:46,236 - distributed.core - INFO - Starting established connection to tls://dask-351d74e2f7f644beb571a17408e81902.staging:8786
2023-03-20 15:14:59,341 - distributed.core - INFO - Event loop was unresponsive in Worker for 9.18s. This is often caused by long-running GIL-holding functions
or moving large chunks of data. This can cause timeouts and instability.
2023-03-20 15:15:00,410 - distributed.utils_perf - INFO - full garbage collection released 47.80 MiB from 0 reference cycles (threshold: 9.54 MiB)
2023-03-20 15:15:03,907 - distributed.core - INFO - Event loop was unresponsive in Worker for 4.53s. This is often caused by long-running GIL-holding functions
or moving large chunks of data. This can cause timeouts and instability.
2023-03-20 15:15:43,799 - distributed.core - INFO - Event loop was unresponsive in Worker for 3.59s. This is often caused by long-running GIL-holding functions
or moving large chunks of data. This can cause timeouts and instability.
2023-03-20 15:15:48,464 - distributed.core - INFO - Event loop was unresponsive in Worker for 4.66s. This is often caused by long-running GIL-holding functions
or moving large chunks of data. This can cause timeouts and instability.
2023-03-20 15:15:51,847 - distributed.core - INFO - Event loop was unresponsive in Worker for 3.38s. This is often caused by long-running GIL-holding functions
or moving large chunks of data. This can cause timeouts and instability.
2023-03-20 15:16:29,763 - distributed.comm.tcp - INFO - Connection from tls://10.244.10.9:44336 closed before handshake completed
2023-03-20 15:18:05,967 - distributed.shuffle._comms - ERROR - Shuffle cc9531b7820cc766da633788babf811c forgotten
Traceback (most recent call last):
  File "/srv/conda/envs/notebook/lib/python3.10/site-packages/distributed/shuffle/_comms.py", line 71, in _process
    await self.send(address, shards)
  File "/srv/conda/envs/notebook/lib/python3.10/site-packages/distributed/shuffle/_worker_extension.py", line 122, in send
    self.raise_if_closed()
  File "/srv/conda/envs/notebook/lib/python3.10/site-packages/distributed/shuffle/_worker_extension.py", line 163, in raise_if_closed
    raise self._exception
RuntimeError: Shuffle cc9531b7820cc766da633788babf811c forgotten
2023-03-20 15:18:05,977 - distributed.shuffle._comms - ERROR - Shuffle cc9531b7820cc766da633788babf811c forgotten
Traceback (most recent call last):
  File "/srv/conda/envs/notebook/lib/python3.10/site-packages/distributed/shuffle/_comms.py", line 71, in _process
    await self.send(address, shards)
```





Logging best practices

Quick review: How to do good logging?

1. Don't use `print()`.
2. Provide the necessary context.
3. Use structured logging.





Distributed logging

How to access all your logs in a distributed environment?

```
>>> cluster.scale(100)
```

Congratulations, you now have 102 different places to look for logs!





Distributed logging

How can I access all my logs in a distributed environment?

Built-in:

```
>>> client.get_scheduler_logs()
```

```
>>> client.get_worker_logs(workers=["10.0.2.96"])
```

Problem: If it's dead, it doesn't return logs.





Centralized logging

How can I access all my logs in a distributed environment?

Better: Centralized logging in an external system

- Persistence
- Querying

```
$ coiled cluster logs --cluster 183522 --workers 10.0.2.96  
--filter "memory usage"
```

```
(10.0.2.96)      2023-04-03 00:28:30.454000 distributed.worker.memory - WARNING - Worker is at 80%  
memory usage. Pausing worker. Process memory: 5.76 GiB -- Worker memory limit: 7.15 GiB  
(10.0.2.96)      2023-04-03 00:28:31.657000 distributed.worker.memory - WARNING - Worker is at 79%  
memory usage. Resuming worker. Process memory: 5.70 GiB -- Worker memory limit: 7.15 GiB
```



The background features a series of thin, light gray vertical lines of varying heights and widths, creating a textured, grid-like effect. A prominent, solid blue vertical bar is positioned to the left of the word 'Metrics', extending from the top of the word down to the bottom of the page.

Metrics



What are metrics?





Metrics of a Dask cluster

`distributed` provides detailed statistics about its parts.

`dask_scheduler_workers`

Number of workers known by scheduler

`dask_worker_memory_bytes`

Memory breakdown per worker

...and many others.

See the full list at <https://distributed.dask.org/en/latest/prometheus.html>.



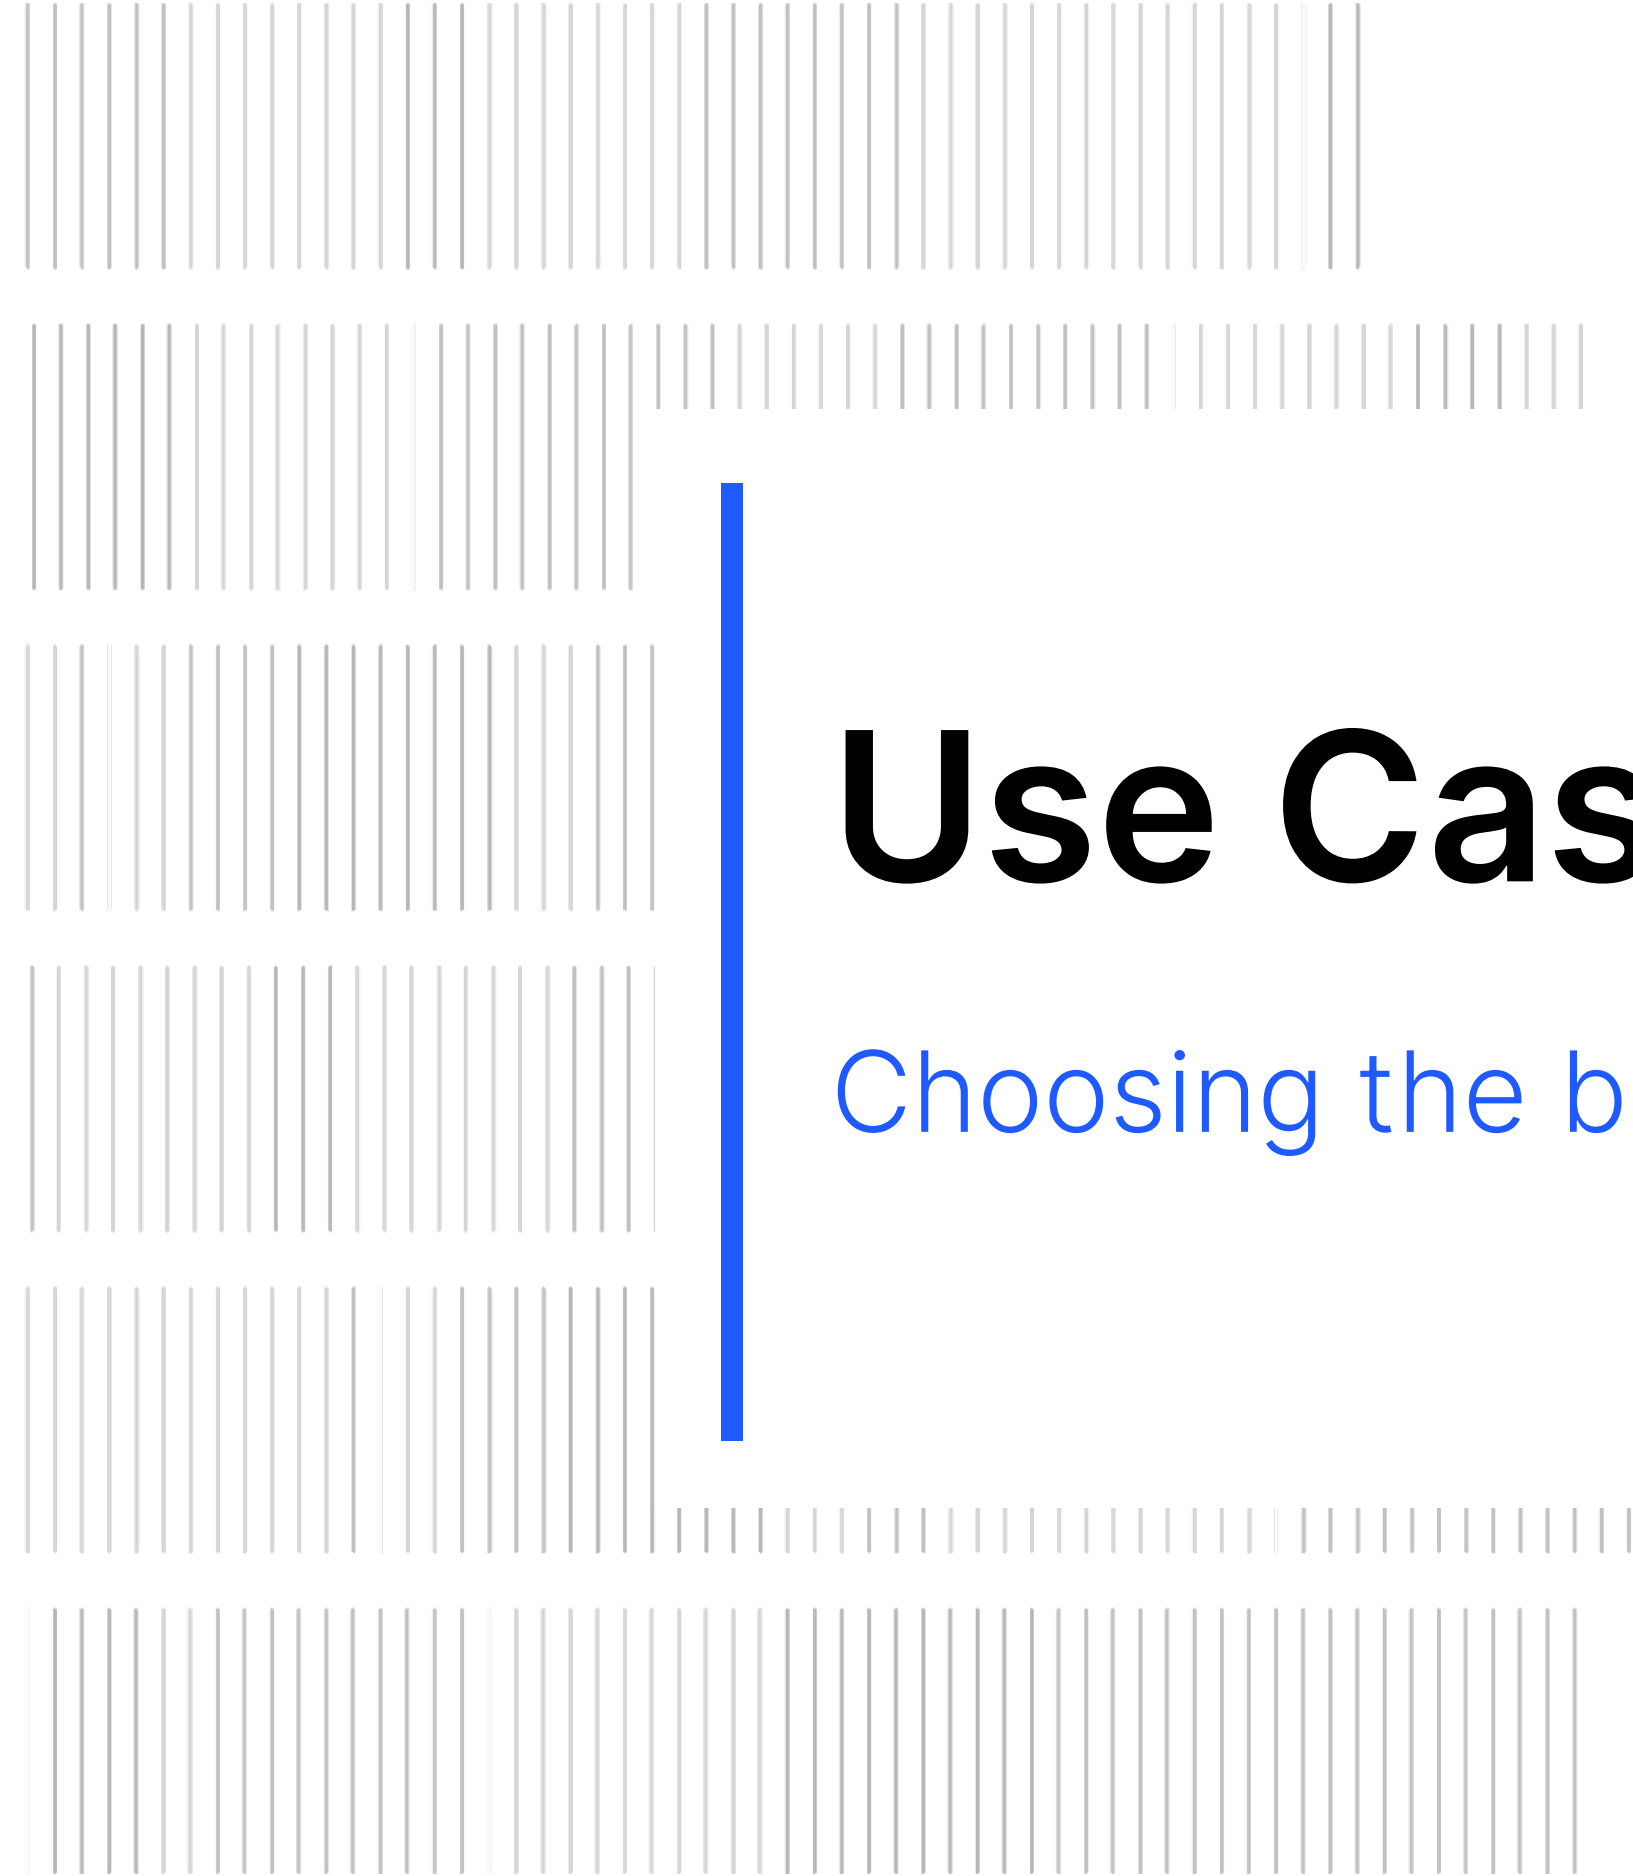


Centralized metrics

Enabling automated monitoring and diagnostics on the entire cluster

- Dask has a native Prometheus integration
- Centralized metrics enable:
 - Powerful dashboards
<https://benchmarks-grafana.oss.coiledhq.com/>
 - Metric-based alerts/flags






Use Case


Choosing the best hardware



Workload 1: Network and disk with array workload



[t3.large](#)




[m6i.large](#)






Workload 2: XGBoost + Optuna with moderate CPU use



[t3.large](#)



[m6i.large](#)





Use case: Choosing the best hardware

Which AWS instance types should we choose as default?

Result: For data science/engineering workloads, non-burstable instances are both cheaper per hour and faster.

Find the full blog post at <https://blog.coiled.io/blog/burstable-vs-nonburstable>.



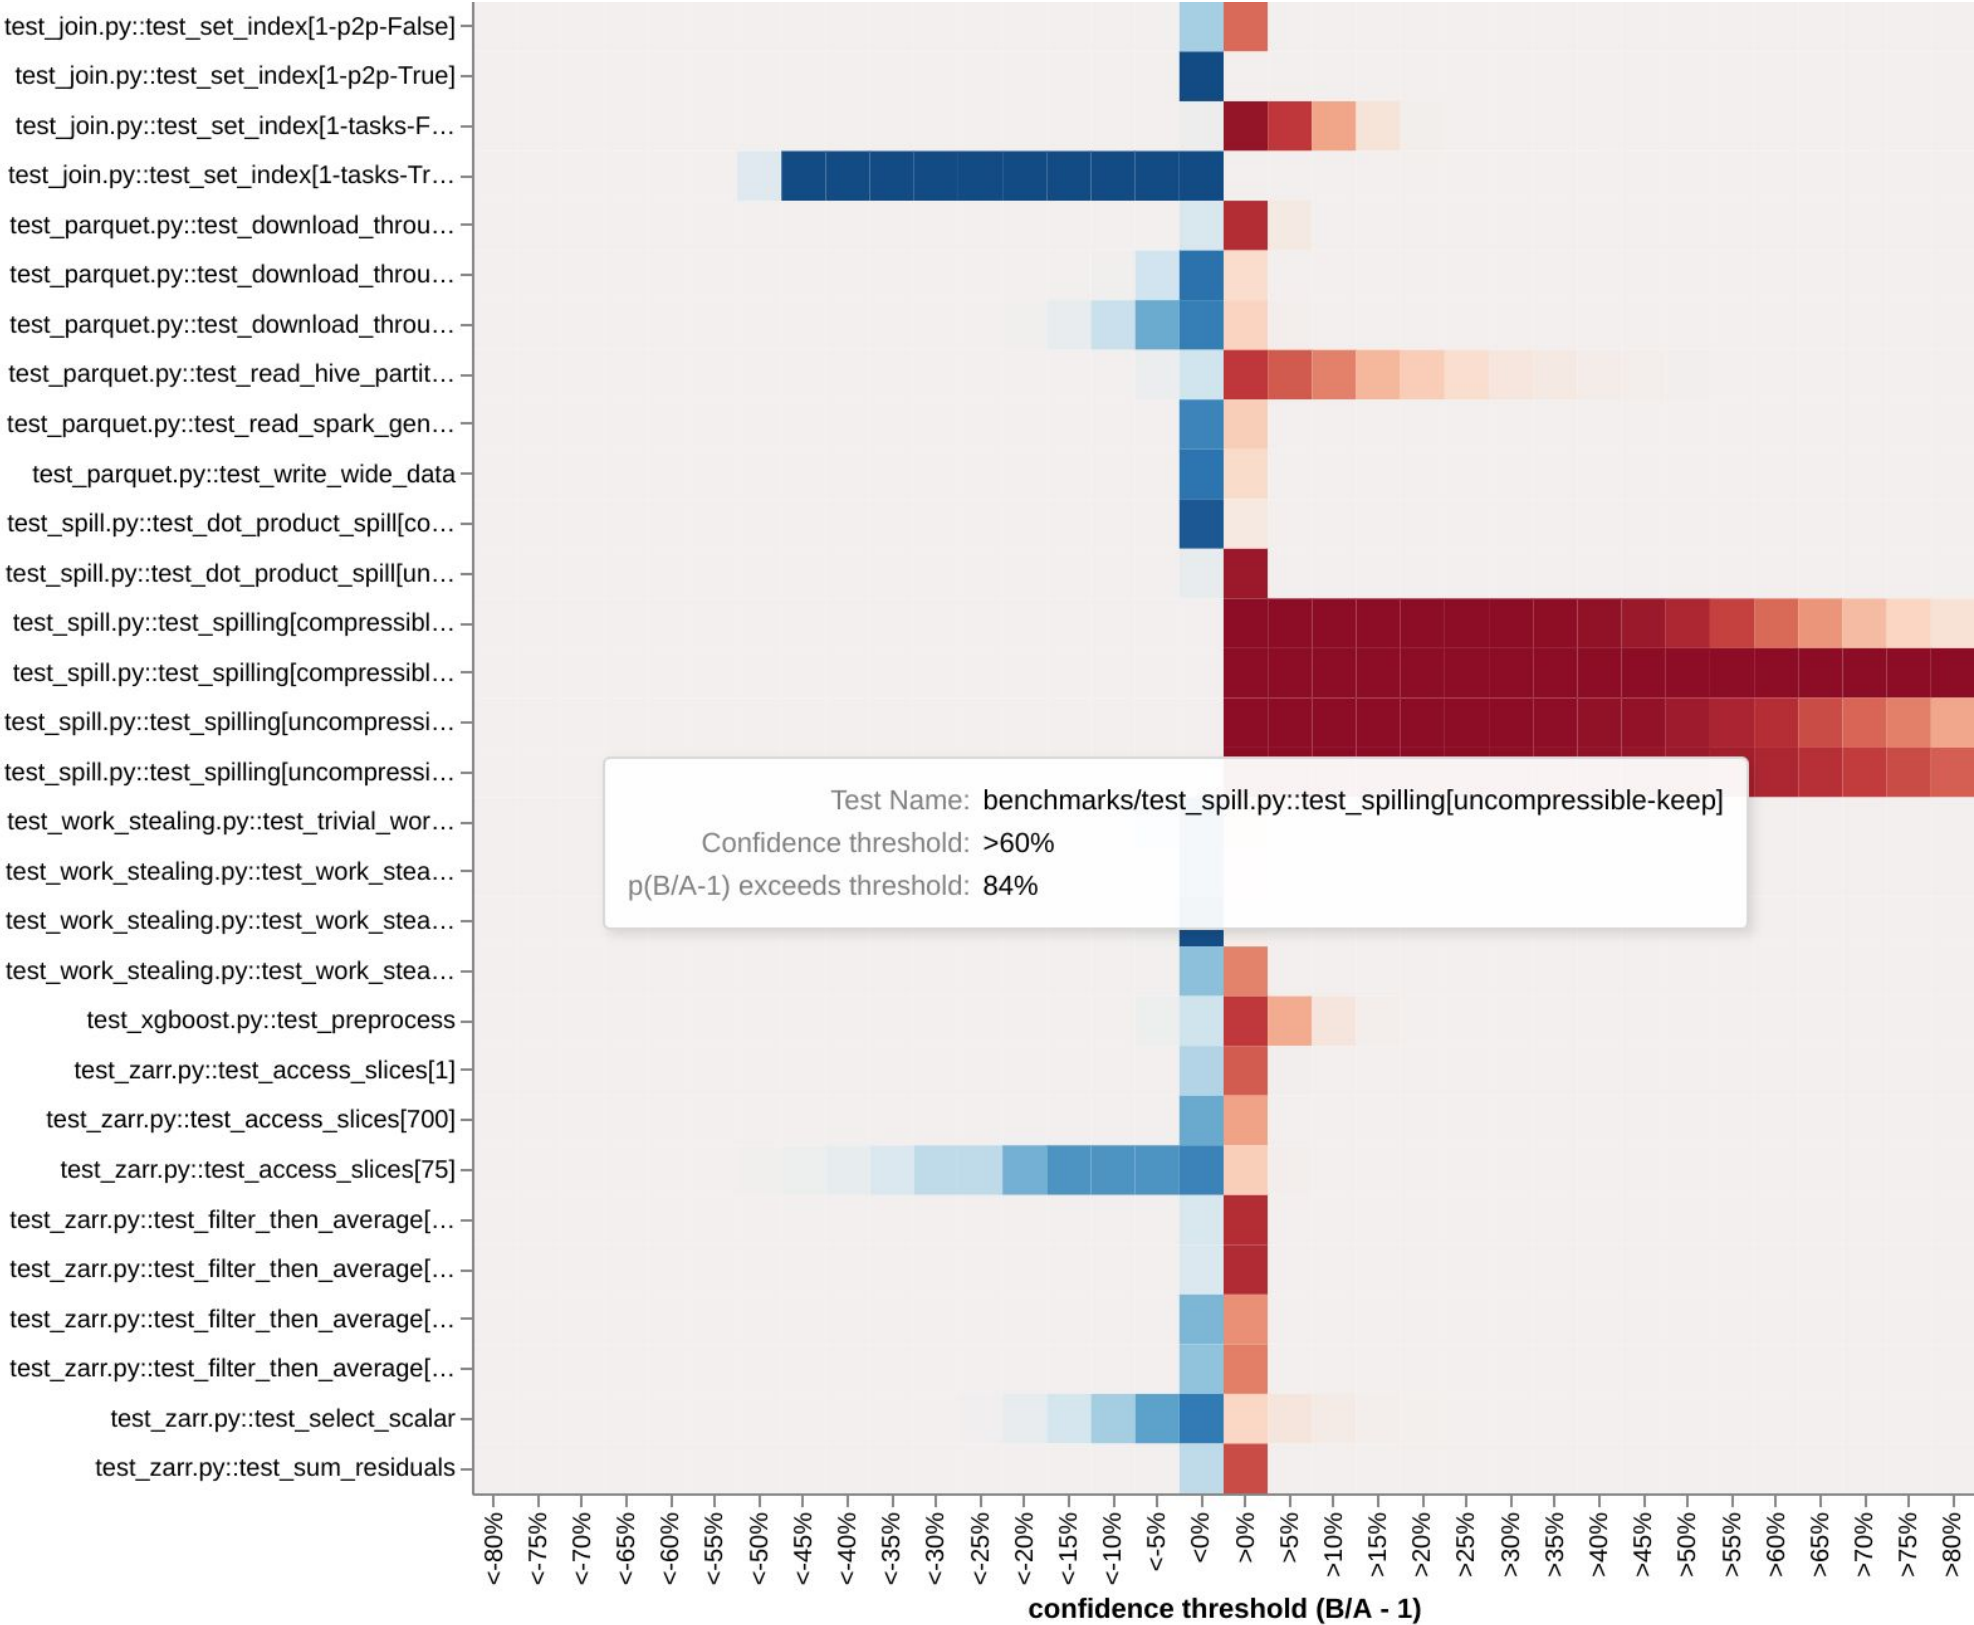


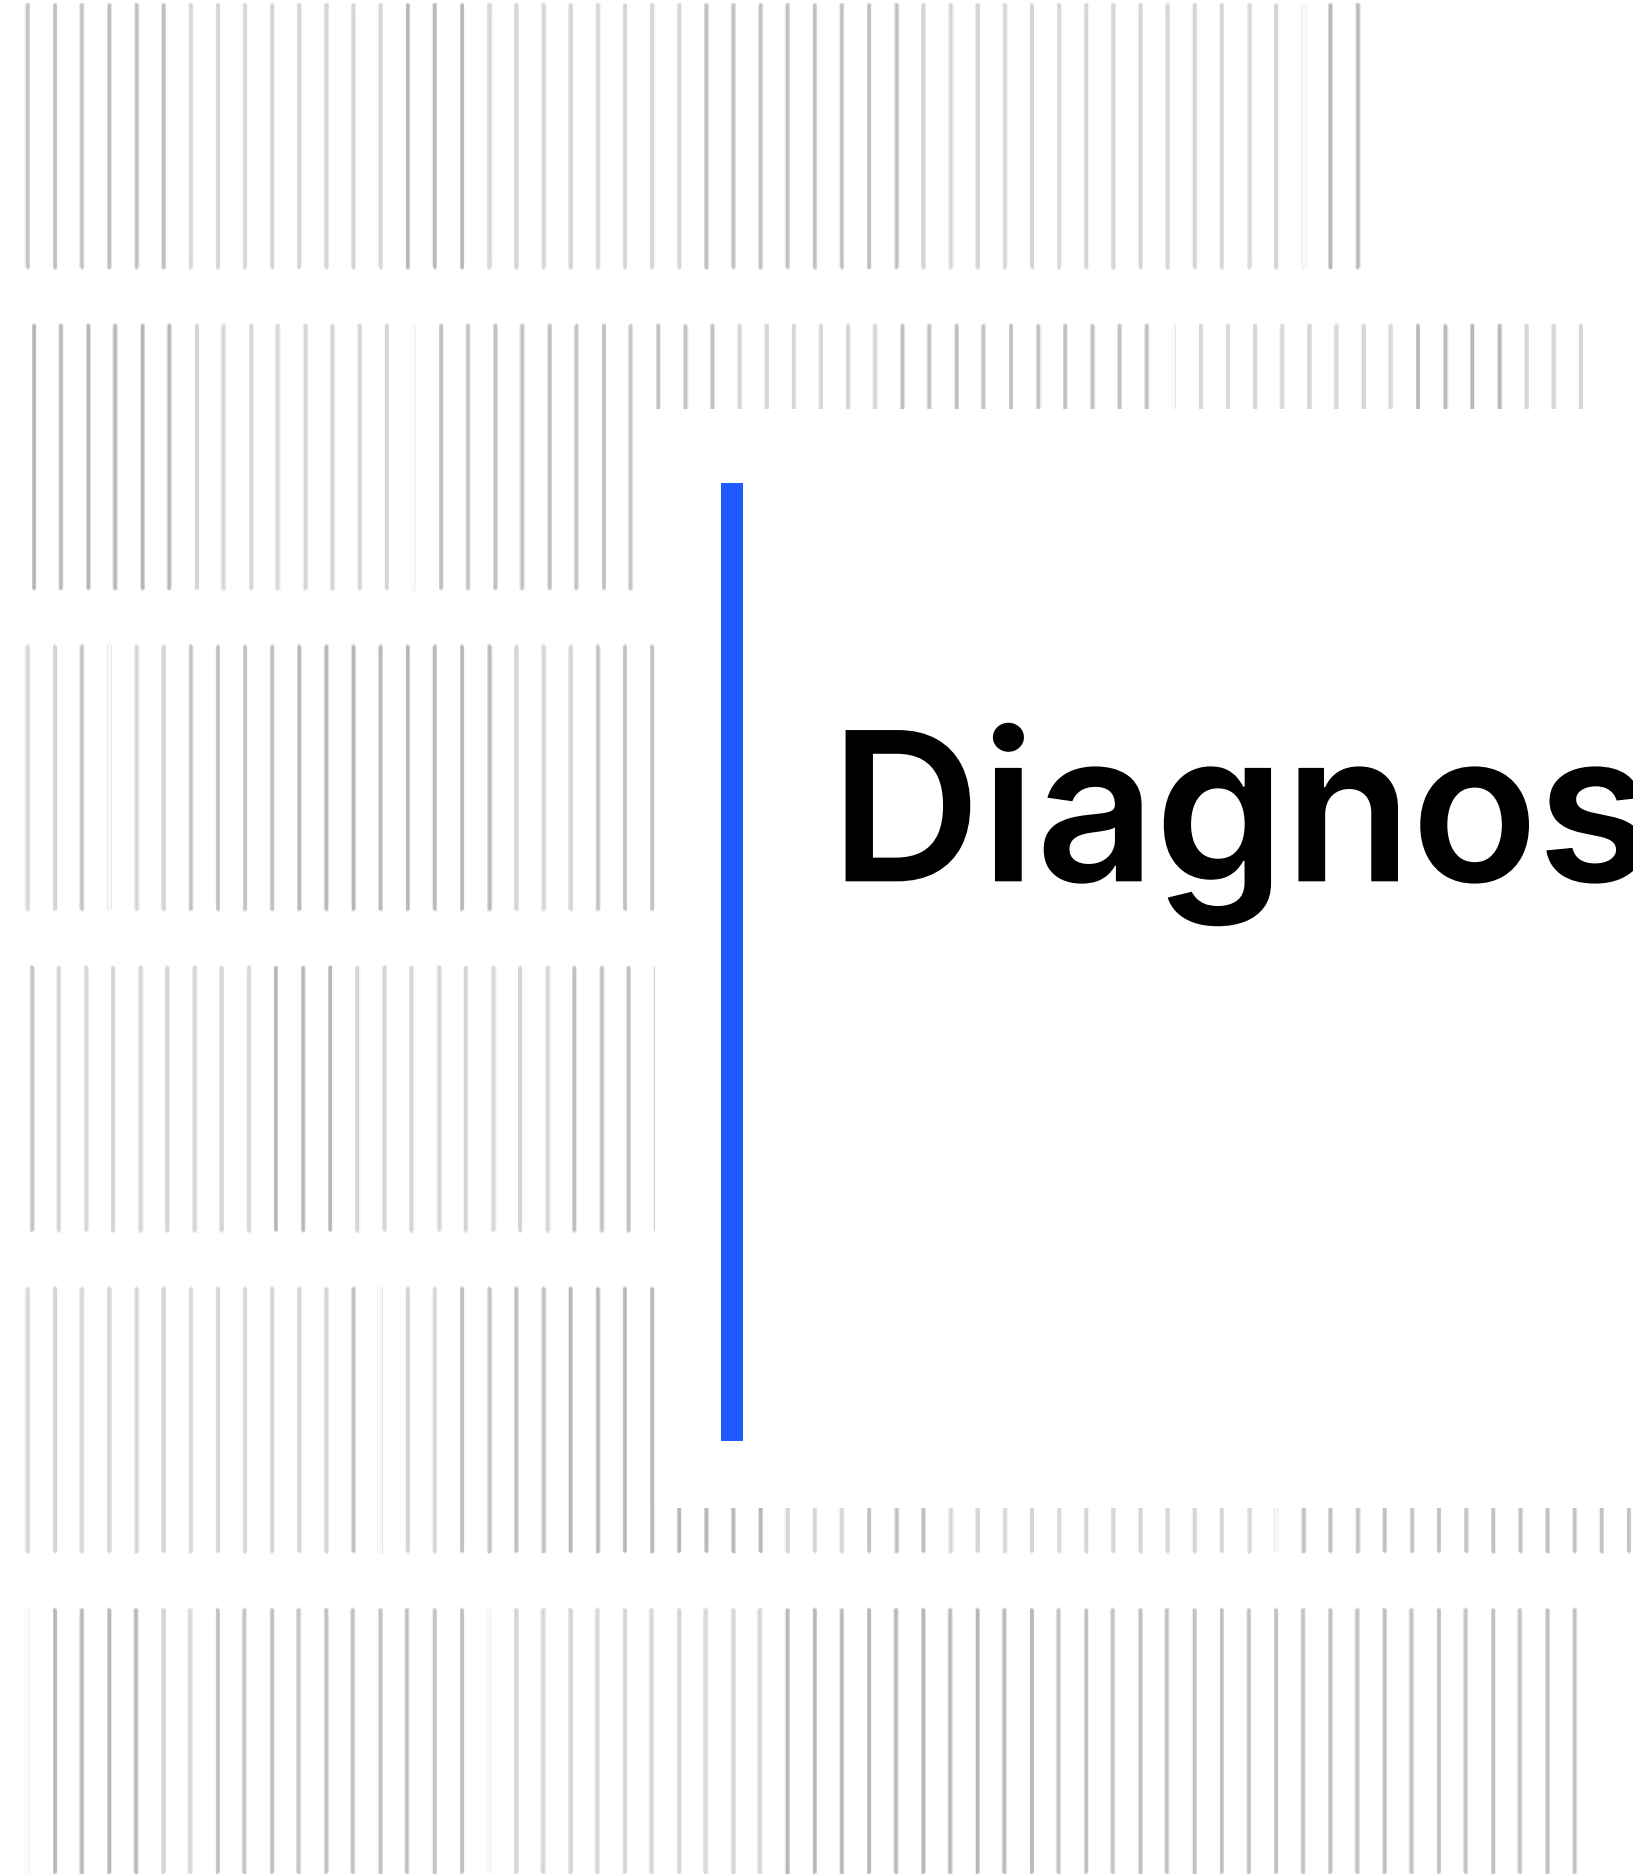
Use case: Data-driven development

Data-driven design for the
Dask scheduler

Learn more at:

**Data-driven design for the
Dask scheduler**
by *Guido Imperiale*
today, 14:10 - 14:40
in Hall B07-B08





Diagnostic Tooling



Dashboards


Tracking your clusters
and your workloads.

Dashboards are

👍 **very useful**

👎 **ephemeral**

...what if we could store
a subset of that data?





Performance reports

A built-in way for storing and sharing performance data about your workload execution

```
from distributed import performance_report

with performance_report(filename="pycon_de.html") :
    ...
```

Let's take a look at an example performance report [here](#).





Profiling

Gaining deep insights into performance

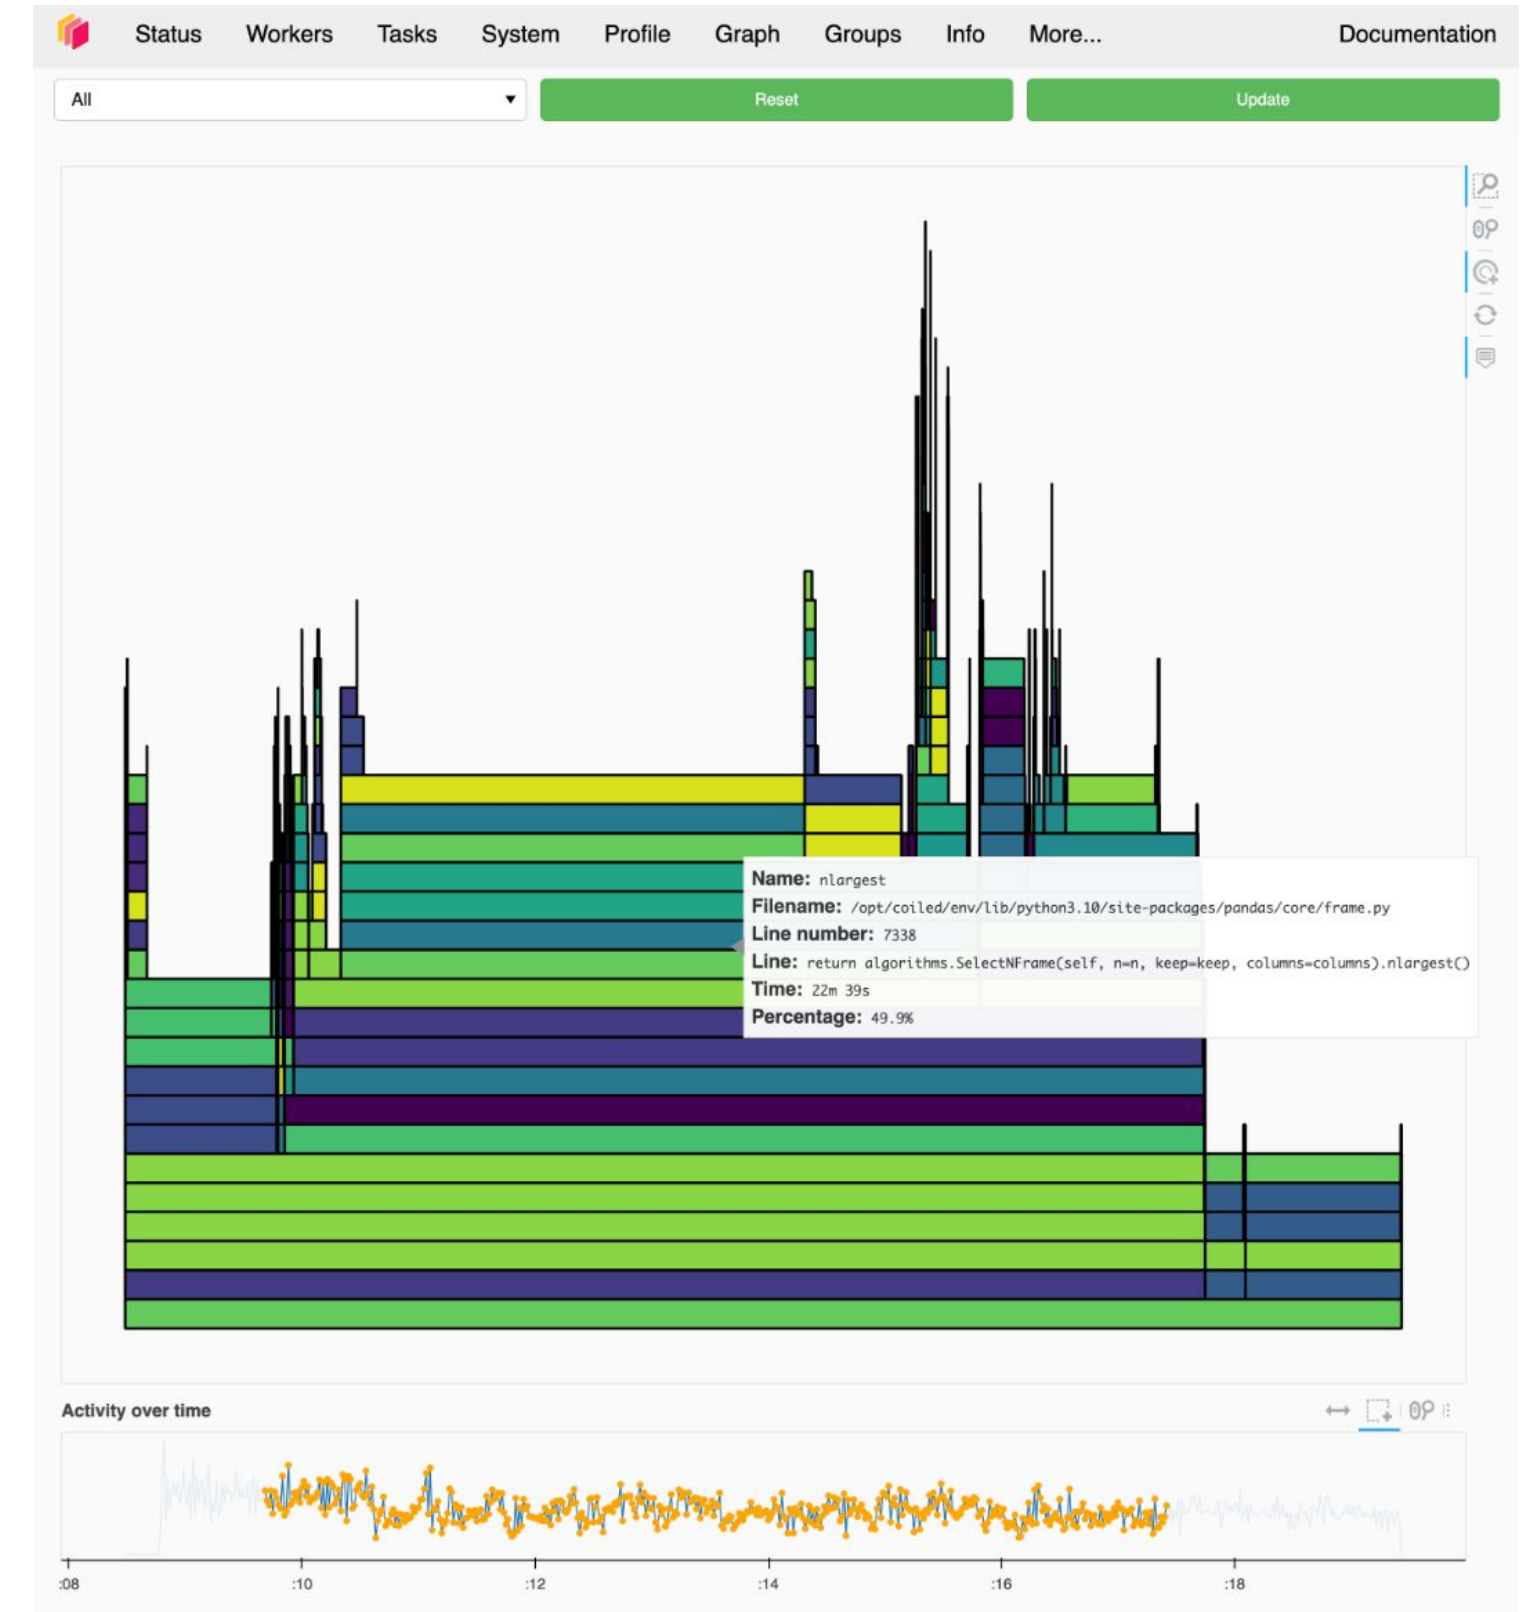
Built-in: **/profile**

- Statistical CPU profiler

- 👍 Low overhead → always-on
- 👍 Cluster aggregates or individual profiles

- 👎 Python-only
- 👎 Lacks rich timeline

For more in-depth profiling: **dask-pyspy**






Zooming back out

- **An observable system lets us ask: Why is this happening?**
- Logging and metrics are key pillars of an observable system.
- Good diagnostic tools refine the telemetry data to create insights.
- Dask has built-in logging, metrics, and diagnostic tools.
- Centrally and durably collect your telemetry data to fully leverage them.





Coiled

A Dask Company

Reliably scale Python in your cloud
account. Get started for free.

coiled.io